

# Hardware-accelerated security monitoring

Marcus Botacin

<sup>1</sup>botacin@tamu.edu  
marcusbotacin.github.io

## Who Am I?

- Assistant Professor (2022) - Texas A&M University (TAMU), USA
  - ACES Program Fellowship
- PhD. in Computer Science (2021) - Federal University of Paraná (UFPR), Brazil
  - Thesis: *“On the Malware Detection Problem: Challenges and new Approaches”*
- MSc. in Computer Science (2017) - University of Campinas (UNICAMP), Brazil
  - Dissertation: *“Hardware-Assisted Malware Analysis”*
- Computer Engineer (2015) - University of Campinas (UNICAMP), Brazil
  - Final Project: *“Malware detection via syscall patterns identification”*

# Topics

- 1 Introduction
  - The Problem
  - Solution
- 2 Signature Matching
  - HEAVEN
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection
  - MINI-ME
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks

# Bottleneck: Real-time monitoring performance penalty




Figure: AV Monitoring Performance.


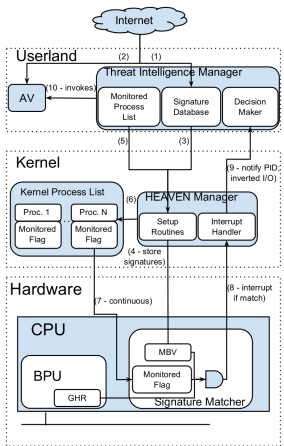


Figure: In-memory AV scans worst-case and best-case performance penalties.

# Topics

- 1 Introduction
  - The Problem
  - **Solution**
- 2 Signature Matching
  - HEAVEN
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection
  - MINI-ME
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks


# Hardware AV Architecture



## 2-level Architecture

Do not fully replace AVs, but add efficient matching capabilities to them.

# Performance Characterization



## 2-Phase HEAVEN CPU Performance

The inspection phase causes occasional, and quick bursts of CPU usage. The AV operating alone incurs a continuous 10% performance overhead.

# Topics

- 1 Introduction
  - The Problem
  - Solution
- 2 **Signature Matching**
  - **HEAVEN**
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection
  - MINI-ME
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks



## Publication



Figure: Source:

<https://www.sciencedirect.com/science/article/abs/pii/S0957417422004882>

# Branch patterns as signatures






Figure 2: **Signature Generation Policy.** Associating high-level code constructs with their occurrence in the execution flow.

# Branch patterns as signatures

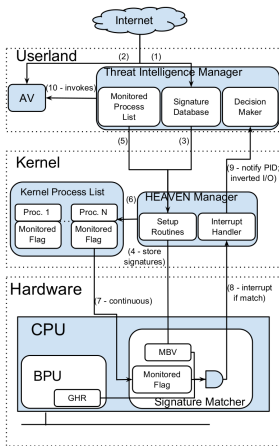


**Figure: Two-level branch predictor.** A sequence window of taken (1) and not-taken (0) branches is stored in the Global History Register (GHR).



**Figure: Branch patterns coverage.**


# Hardware AV Architecture



## 2-level Architecture

Do not fully replace AVs, but add efficient matching capabilities to them.

# Performance Characterization



## 2-Phase HEAVEN CPU Performance

The inspection phase causes occasional, and quick bursts of CPU usage. The AV operating alone incurs a continuous 10% performance overhead.

# Topics

- 1 Introduction
  - The Problem
  - Solution
- 2 Signature Matching
  - HEAVEN
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection
  - MINI-ME
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks

# Publication

## The AV says: Your Hardware Definitions Were Updated!

Publisher: IEEE


[Cite This](#)

[PDF](#)


Marcus Botacin ; Lucas Galante ; Fabricio Ceschin ; Paulo C. Santos ; Luigi Carro ; Paulo de Geus ; André Grégio ; Marco A. Z. ... [All Authors](#)

Figure: Source: <https://ieeexplore.ieee.org/document/9034972/>

# Profiling-Based AV



**Figure: Malware Classification using low level features.**



**Figure: REHAB Architecture.** CPU's HPC data is used as feature for a FPGA-based, reconfigurable ML classifier updatable via software.



# Classifiers




Figure: SVM.




Figure: Random Forest.




Figure: MLP.

## AV Checks Cost

**Table: Execution Speedup per AV check.** Hardware Accelerator is essential for overhead elimination.

| ML algorithm →   | SVM         | RF          | MLP         |
|------------------|-------------|-------------|-------------|
| <b>CPU</b>       | 220 $\mu$ s | 270 $\mu$ s | 240 $\mu$ s |
| <b>FPGA+Comm</b> | 124.5ns     | 111.2ns     | 158.9ns     |
| <b>Speedup</b>   | 1.7k×       | 2.4k×       | 1.5k×       |

# Topics

- 1 Introduction
  - The Problem
  - Solution
- 2 Signature Matching
  - HEAVEN
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection
  - MINI-ME
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks

## Publication

Original Paper | [Published: 13 February 2020](#)

### The self modifying code (SMC)-aware processor (SAP): a security look on architectural impact and support


[Marcus Botacin](#) ✉, [Marco Zanata](#) & [André Grégio](#)

*Journal of Computer Virology and Hacking Techniques* **16**, 185–196(2020) | [Cite this article](#)

**198** Accesses | **3** Altmetric | [Metrics](#)

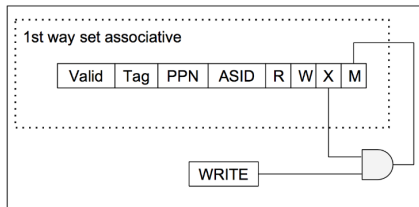
**Figure:** **Source:** <https://link.springer.com/article/10.1007/s11416-020-00348-w>

# Architectural Support



**Fig. 2** Effectiveness of event counter as a SMC detector.

Figure: Pipeline Stalls Detection.



**Fig. 4** MMU-based SMC detection mechanism.

Figure: MMU Modification.

# Page Handling Overhead

```
1 static noinline void __do_page_fault(...)
  {
2 // Original Code
3 if(kprobes())...
4 // Instrumentation Code
5 if(was_executable_page_written()) {
6     if(!is_allowed_process(get_pid())) {
7         // SMC Detected
```

**Code 9** SMC detection routines in the Linux kernel. The added verification instructions are executed every page fault.

Figure: Page Fault Handler.

**Table 3** Estimated overhead of Software-based SMC detectors during page fault trapping on SPEC applications

| Benchmark | Penalty | Benchmark | Penalty | Benchmark  | Penalty |
|-----------|---------|-----------|---------|------------|---------|
| bzip2     | 3.47%   | mcf       | 3.76%   | wrf        | 3.91%   |
| namd      | 1.54%   | bwaves    | 3.73%   | perlbench  | 6.49%   |
| h265ref   | 4.61%   | calculix  | 3.57%   | dealli     | 3.12%   |
| astar     | 2.12%   | sjeng     | 2.74%   | hammer     | 2.62%   |
| gobmk     | 3.10%   | cactusADM | 3.70%   | libquantum | 3.24%   |
| gcc       | 6.25%   | gromacs   | 4.01%   | sphinx3    | 3.76%   |
| lbm       | 4.27%   | zeusmp    | 3.48%   | povray     | 4.64%   |
| tonto     | 4.53%   | GemsFDTD  | 3.48%   | xalancbmk  | 3.85%   |
| gamess    | 4.05%   | leslie3d  | 3.46%   | specrand   | 3.36%   |

Figure: Performance Penalty.

# Topics




- 1 Introduction
  - The Problem
  - Solution
- 2 Signature Matching
  - HEAVEN
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection**
  - MINI-ME**
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks

# Publication

RESEARCH-ARTICLE

## Near-Memory & In-Memory Detection of Fileless Malware



**Authors:**  [Marcus Botacin](#),  [André Grégio](#),  [Marco Antonio Zanata Alves](#) [Authors Info & Affiliations](#)

**Publication:** MEMSYS 2020: The International Symposium on Memory Systems • September 2020 • Pages 23–38 • <https://doi.org/10.1145/3422575.3422775>

**Figure:** **Source:** <https://dl.acm.org/doi/10.1145/3422575.3422775>



## MINI-ME

# Malware Identification based on Near- and In-Memory Evaluation (MINIME)

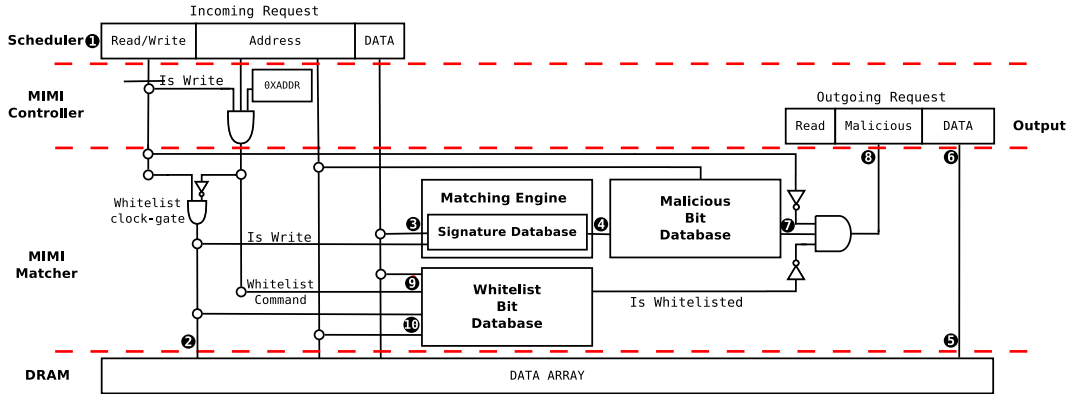



Figure: MINIME Architecture.

# Performance Gains

## MINIME vs. On-Access AVs

Significant performance gains even in the worst case.



# Topics

- 1 Introduction
  - The Problem
  - Solution
- 2 Signature Matching
  - HEAVEN
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection
  - MINI-ME
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks

# Publication

## TERMINATOR: A Secure Coprocessor to Accelerate Real-Time AntiViruses using Inspection Breakpoints



Marcus Botacin, Federal University of Paraná (UFPR-BR)  
Francis B. Moreira, Federal University of Rio Grande do Sul (UFRGS-BR)  
Philippe O. A. Navaux, Federal University of Rio Grande do Sul (UFRGS-BR)  
André Grégio, Federal University of Paraná (UFPR-BR)  
Marco A. Z. Alves, Federal University of Paraná (UFPR-BR)

Figure: Source: <https://dl.acm.org/doi/10.1145/3494535>

# Function Arguments Inspection




Figure 1: **Function Interposition.** A trampoline function added by AVs to interpose the original function calls.

Figure: Function Interposition.

```

1 rule IsPacked : PECheck {
2   condition:
3     // MZ signature at offset 0 and
4     uint16(0) == 0x5A4D and
5     // PE signature at offset stored
6     // in MZ header at 0x3C
7     uint32(uint32(0x3C)) == 0x00004550
8     and
9     math.entropy(0, filesize) >= 7.0
10 }
  
```

Code 1: YARA rule to detect packed PE files.

Figure: Matching Framework.

# Inspection Triggering

```
PcreateProcessNotifyRoutine
(...){
    pid = GetProcessId();
    libs = EnumProcessModules(
        pid);
    addr = GetModuleAddress(
        libs[target_lib],
        target_function);
    VirtualProtect(addr,
        WRITABLE);
    __intrinsics_set_trap(addr);
    VirtualProtect(addr,
        NOT_WRITABLE|EXECUTABLE)
    ;
}
```

Code 2: Process Creation Callback.

```
1 target_function(...)
2 {
3     entry_checkpoint(
4         empty_slot()
5     );
6     first_task();
7     internal_calls();
8     ...
9     exit_checkpoint();
10    return;
11 }
```

Code 3: Target Function Code.

**Figure: Inspection Breakpoint.**

# Parallel Execution Constraints





Figure 10: **Weighted Average.** Extreme values are unlikely to be monitored and thus were discarded.

Figure: **Scanning Cycles Boundary.**

# Performance Penalty Reduction



**Figure: Performance evaluation when tracking all function calls.** Comparison between execution without AV (BASE), execution with software AV, and execution with the proposed coprocessor model.



# Topics

- 1 Introduction
  - The Problem
  - Solution
- 2 Signature Matching
  - HEAVEN
- 3 HPC Classification
  - REHAB
- 4 Packer Identification
  - SAP
- 5 Fileless Malware Detection
  - MINI-ME
- 6 Function Checking
  - TERMINATOR
- 7 Conclusions
  - Recap & Remarks

# Summary

## Malware Detection

- Huge performance penalties.
- Increasing performance increases detection.

## Academic Contributions

- Branch patterns to replace byte-based signatures.
- FPGAs to classify HPCs in runtime.
- SMC-aware processor to detect packers.
- Instrumented memory controller to detect fileless malware.
- CPU coprocessors for real-time syscall checking.

# Thanks!

Questions? Comments?

@MarcusBotacin  
botacin@tamu.edu  
marcusbotacin.github.io